

AFSAL A AZEEZ

Full-Stack & AI Engineer • GenAI • IIT Bombay M.Tech

afsalaazeez@gmail.com • +971 55 388 5320

Abu Dhabi, UAE • **Available Immediately**

linkedin.com/in/afsalaazeez • github.com/zencodelab • afsal-a-azeez.in

Professional Summary

Full-Stack & AI Engineer with a profound technical background specializing in designing, building, and scaling robust distributed microservices and production-grade AI/LLM applications. Proven track record at a Silicon Valley venture-backed startup driving technical strategy, optimizing backend APIs, and engineering multi-LLM orchestration platforms. Recently architected and shipped Markaba AI, an AI-powered automotive diagnostics and marketplace platform, solo from first commit to live production. Expert in cloud-native infrastructures (AWS/Kubernetes) and high-performance system design. Graduated from IIT Bombay (GATE AIR-241); fully attested academic credentials for immediate UAE residency compliance.

Technical Skills

- **Architecture & Design:** Microservices, Distributed Systems, RESTful/GraphQL APIs, Serverless, Event-Driven Architecture, System Design.
- **Frontend Engineering:** React.js, Next.js, Electron.js, Redux Toolkit, Zustand, TailwindCSS, SCSS, HTML5/CSS3.
- **Backend Engineering:** Node.js (TypeScript), Python (FastAPI/Django), Go, Express, Fastify, OAuth, JWT.
- **AI & LLM:** GPT-4, Gemini, Claude, Llama 3, LangChain, OpenAI API, RAG Pipelines, Vector Databases, Multi-LLM Orchestration, AI Agent Orchestration.
- **DevOps & Cloud:** AWS (EC2, S3, Lambda, EKS), Docker, Kubernetes, OpenShift, Terraform, Helm, CloudFormation, GitHub Actions, Jenkins.
- **Databases & Caching:** PostgreSQL, MySQL, MongoDB, Redis, SQLite, Elasticsearch.
- **Testing & Security:** Playwright, OWASP ZAP, Postman, Artillery, Selenium, JUnit, pytest, Jest.
- **Version Control & Platforms:** GitHub, GitLab, Bitbucket, Azure Repos, Microsoft Teams (Bot Framework).

Professional Experience

Founder & Lead Engineer

Jul 2026 – Present

Full-Stack & AI Engineer

Markaba AI (markabaai.com) — AI Automotive Diagnostics Platform

- **AI Diagnosis Engine:** Architected a multi-layer automotive intelligence engine (Diagnose → Match → Rank → Quote) powered by the Anthropic Claude API, converting natural-language vehicle symptoms into ranked repair plans with matched parts and bookable workshop quotations.
- **Semantic Parts Matching:** Built a RAG-based parts-matching pipeline using LlamaIndex and Ollama embeddings over a PostgreSQL + pgvector catalog, with keyword-matching fallback for reliability.
- **Full-Stack Marketplace:** Shipped a bilingual (English/Arabic, RTL) React + Vite + TypeScript + Tailwind frontend and a FastAPI + SQLAlchemy + Alembic backend covering a parts marketplace, vehicle-scoped “Mod Discovery” upgrades, workshop/service booking, and a “My Garage” vehicle-profile system.
- **Payments & Notifications:** Built a Tap Payments integration, Clerk/OIDC authentication, and SendGrid/Twilio (email, SMS, WhatsApp) workflows, including a WhatsApp Accept/Decline flow for partner workshops.
- **Solo Production Delivery:** Owned architecture, backend, frontend, and DevOps end-to-end — Dockerized services behind Caddy on a VPS with GitHub Actions CI/CD — taking the product from zero to a live production deployment within three weeks of the first commit.
- **Tech Stack:** Python, FastAPI, React, Vite, TypeScript, Tailwind CSS, PostgreSQL, pgvector, LlamaIndex, Docker, GitHub Actions, Anthropic Claude API.

Member of Technical Staff

Sep 2020 – May 2026

Full-Stack & AI Engineer

Zettabytes Inc. (Roost.ai) — Mayfield-funded DevOps & AI Startup · Remote /

US-Based Market

Phase B — RoostGPT: AI Test Generation Platform

2023 – 2026

- **RoostGPT — AI Testing Platform:** Built and shipped **RoostGPT**, an AI testing copilot leveraging LLMs (GPT-4, Gemini, Claude, Llama 3) to auto-generate unit, API, UI, and integration test cases from source code, Swagger/OpenAPI specs, and Jira user stories.
- **Microsoft Teams Notification Bot:** Took **end-to-end ownership** of building a Microsoft Teams notification bot for RoostGPT, delivering real-time test generation status updates, results summaries, and CI/CD pipeline notifications directly within Teams channels.

- **Git Platform Integrations:** Built integrations within RoostGPT connecting to GitHub, GitLab, Bitbucket, and Azure Repos for automated repository access, test generation triggers, and direct test commit workflows. Implemented GitHub pull request webhook hooks to automatically retrigger test generation on new code changes.
- **Multi-LLM Provider UI:** Built the frontend dashboard for configuring and switching between LLM providers — supporting OpenAI, Google Gemini, Anthropic Claude, and Meta Llama 3 as interchangeable backends, enabling customers to select their preferred AI model for test generation.
- **Java Unit Test Generation — End-to-End Ownership:** Owned the complete Java unit test generation pipeline — from repository ingestion through to final generated test files. Engineered LLM context by working directly on the **Spring Boot Java AST parsing service** to extract and structure class graphs, bean dependency trees, and method signatures from multi-module Maven/Gradle repositories. Diagnosed and resolved prompt engineering issues causing incorrect imports and missing mocks, achieving a **93% compile success rate** across generated test suites.
- **Test Scenario Analysis & Report UI:** Built the end-to-end test scenario analysis feature (frontend, API, and backend) — listing all generated tests with covered and uncovered scenarios, allowing users to select specific scenarios and trigger targeted test generation. Displayed generation summaries, compile check results, and parsed coverage report output, giving developers a full view of test quality before review and commit. Achieved **4x improvement in test coverage** and **85% reduction in manual testing time**.
- **Prompt Management System:** Designed and shipped a **nested-map data structure** for managing LLM prompts across multiple dimensions — model (GPT-4, Gemini, Claude, Llama 3), test type (unit, API spec, integration), and client. Implemented a deterministic lookup algorithm enabling the system to fetch the exact prompt for any combination at runtime, and enforced structured folder conventions for prompt files with **AWS S3-backed storage** for versioning and distribution.
- **Java AST Schema Redesign:** Flattened nested class structures into first-class entities, reducing per-class file size from **316KB to 17KB** (95% reduction) — improving LLM context efficiency and test generation accuracy.
- **FQN Mapping & Method Overload Resolution:** Implemented Fully Qualified Name mapping for internal method calls, enabling accurate method overload identification via argument lists.
- **File-to-Class Entity Migration:** Led migration from file-based to class-based entity handling across the test generation pipeline, enabling AI agents to correctly receive structured parameters for Java unit test generation.
- **AST Output Optimisation:** Removed unnecessary body fields from method call extraction, reducing codeparse output from **15.9MB to significantly smaller payloads** — cutting LLM token costs and improving throughput.
- **Schema Validation at Scale:** Validated AST parser output against multiple Java repositories (ask-java, adk-java) to confirm schema correctness at enterprise scale.
- **Test Generation UI:** Built the full frontend for test generation workflows — repository analysis views, trigger dashboards, and review/refinement interfaces for inspecting, editing, and approving AI-generated tests.
- **User Access Management UI:** Built frontend for team-level permission controls over test generation workflows across repositories and projects.
- **AI Response Caching:** Designed an SQLite-based caching layer for LLM responses to eliminate redundant API calls, reducing cost and latency for repeated prompts.
- **Playwright E2E Testing:** Wrote and reviewed Playwright end-to-end test suites and built an API test analysis page for scenario-specific execution tracking.
- **Technical Mentorship:** Mentored junior engineers and interns; conducted code reviews and tracked delivery via JIRA.
- **Tech Stack:** React.js, Next.js, TypeScript, Node.js, Spring Boot (Java), SQLite, REST APIs, LLM APIs (OpenAI, Gemini, Claude, Llama 3), Docker, GitHub Actions.

Phase A — Roost Desktop, Serverless & EaaS Platform

2020 – 2022

- **Roost Desktop Application:** Built and maintained a **cross-platform Electron.js + React.js** desktop application for DevOps collaboration targeting macOS, Windows, and Linux.
- **Kubernetes Visualisation:** Developed frontend UI for visualising local Kubernetes deployment environments — pod network topology, sidecar/service-mesh connections, and real-time cluster state monitoring.
- **Serverless Migration:** Transitioned the Roost Desktop product to a **serverless cloud architecture**, building the Roost Web UI as a browser-based interface extending the desktop app's capabilities.
- **Collaboration Features:** Implemented a collaboration layer enabling developers to share environment snapshots and deployment states with QA, product, and external stakeholders for review.
- **EaaS Feature Development:** Implemented multi-repo commit (Save All), repo-level config saving, environment variable injection, Helm chart filtering, CloudFormation (S3/JSON) integration, and Terraform UI with full API integration.
- **GitHub API Integration:** Built integrations with GitHub API for repository file discovery, deployment config parsing, and Kubernetes/Helm YAML filtering for the EaaS config workflow.
- **Slack & JIRA Config UI:** Added Slack token/channel suggestion dropdowns and JIRA-referenced feature track-

ing fields across the EaaS configuration UI.

- **Enterprise Deployments:** Supported on-premise and OpenShift deployments for enterprise clients across the US market, collaborating directly with their engineering and DevOps teams on cluster integration and platform rollout.
- **Tech Stack:** React.js, Node.js, Python, Go, REST/GraphQL APIs, Kubernetes, Docker, Terraform, AWS (EKS, EC2, S3, Lambda), Electron.js.

Education

Master of Technology (M.Tech.) in Computer Science & Engineering **Graduated 2020**
Indian Institute of Technology (IIT), Bombay *Mumbai, India*

- GATE AIR-241 — Top 0.22% out of 107,893 candidates. Degree and academic credentials fully attested for UAE international visa and residency compliance.

Bachelor of Technology (B.Tech.) in Computer Science & Engineering **Graduated 2017**
Mar Athanasius College of Engineering (MACE) *Kerala, India*

- Completed B.Tech in CSE with high academic standing.

Technical, AI & Academic Projects

VisionLog-MLX **Personal / Open-Source**

- A MLX-VLM powered vision logging and analysis system utilizing the Qwen2-VL-2B-Instruct model to analyze camera feeds and capture events.

GIS-Integrated Water Network Database System **M.Tech Project – IIT Bombay**

- Designed and developed a GIS-integrated water network database application in QGIS using Python and SQL. Programmed custom Python plugins to automate node elevation extraction from contour files and web APIs, and built efficient data extraction workflows for EPANET and JalTantra.

Slack-like Real-Time Messaging Platform **Academic Project – IIT Bombay**

- Architected a Flask-based real-time coordination and workspace system using Socket.IO for bi-directional browser-to-server messaging. Integrated Elasticsearch for lightning-fast message queries and persistent document indexing.

Flight Delay Prediction Engine **Academic Project – IIT Bombay**

- Built a binary classification system using Python and XGBoost to predict delays exceeding 15 minutes. Applied advanced feature engineering and SMOTE techniques to combat class imbalance, achieving an AUC-ROC of 0.74.

Teaching Assistantships

Graduate Teaching Assistant **Jul 2018 – Nov 2019**
Indian Institute of Technology (IIT), Bombay *Mumbai, India*

- **Software Lab (CS699):** Guided M.Tech and Ph.D. students in programming assignments, managed lab sessions, and performed continuous evaluation and grading of lab work under Prof. Kavi Arya.
- **Computer Programming (CS101):** Led lab sections for 400+ undergraduate students in a team of 45+ TAs under Prof. Om P. Damani and Prof. Ganesh Ramakrishnan; mentored 14 undergraduate students and evaluated exams.